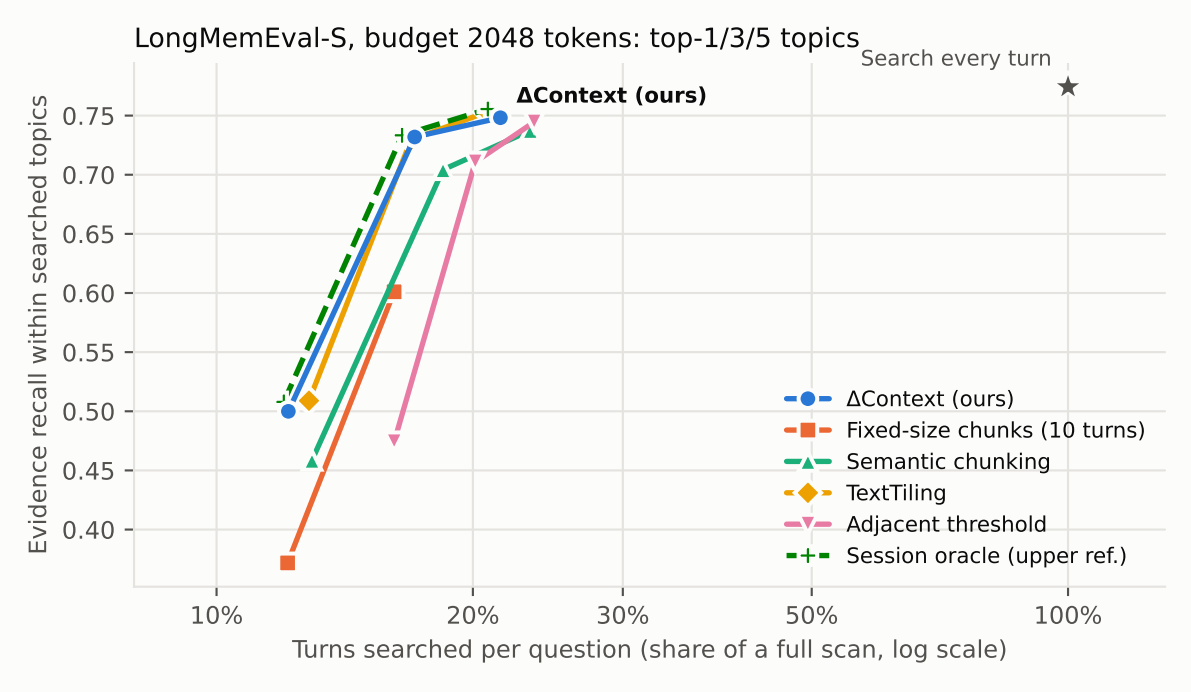
LongMemEval-S, budget 2048 tokens: top-1/3/5 topics
Search every turn
ΔContext (ours)
Evidence recall within searched topics
0.75
0.70
0.65
0.60
0.55
0.50
0.45
0.40
10%
20%
30%
50%
100%
Turns searched per question (share of a full scan, log scale)
ΔContext (ours)
Fixed-size chunks (10 turns)
Semantic chunking
TextTiling
Adjacent threshold
Session oracle (upper ref.)