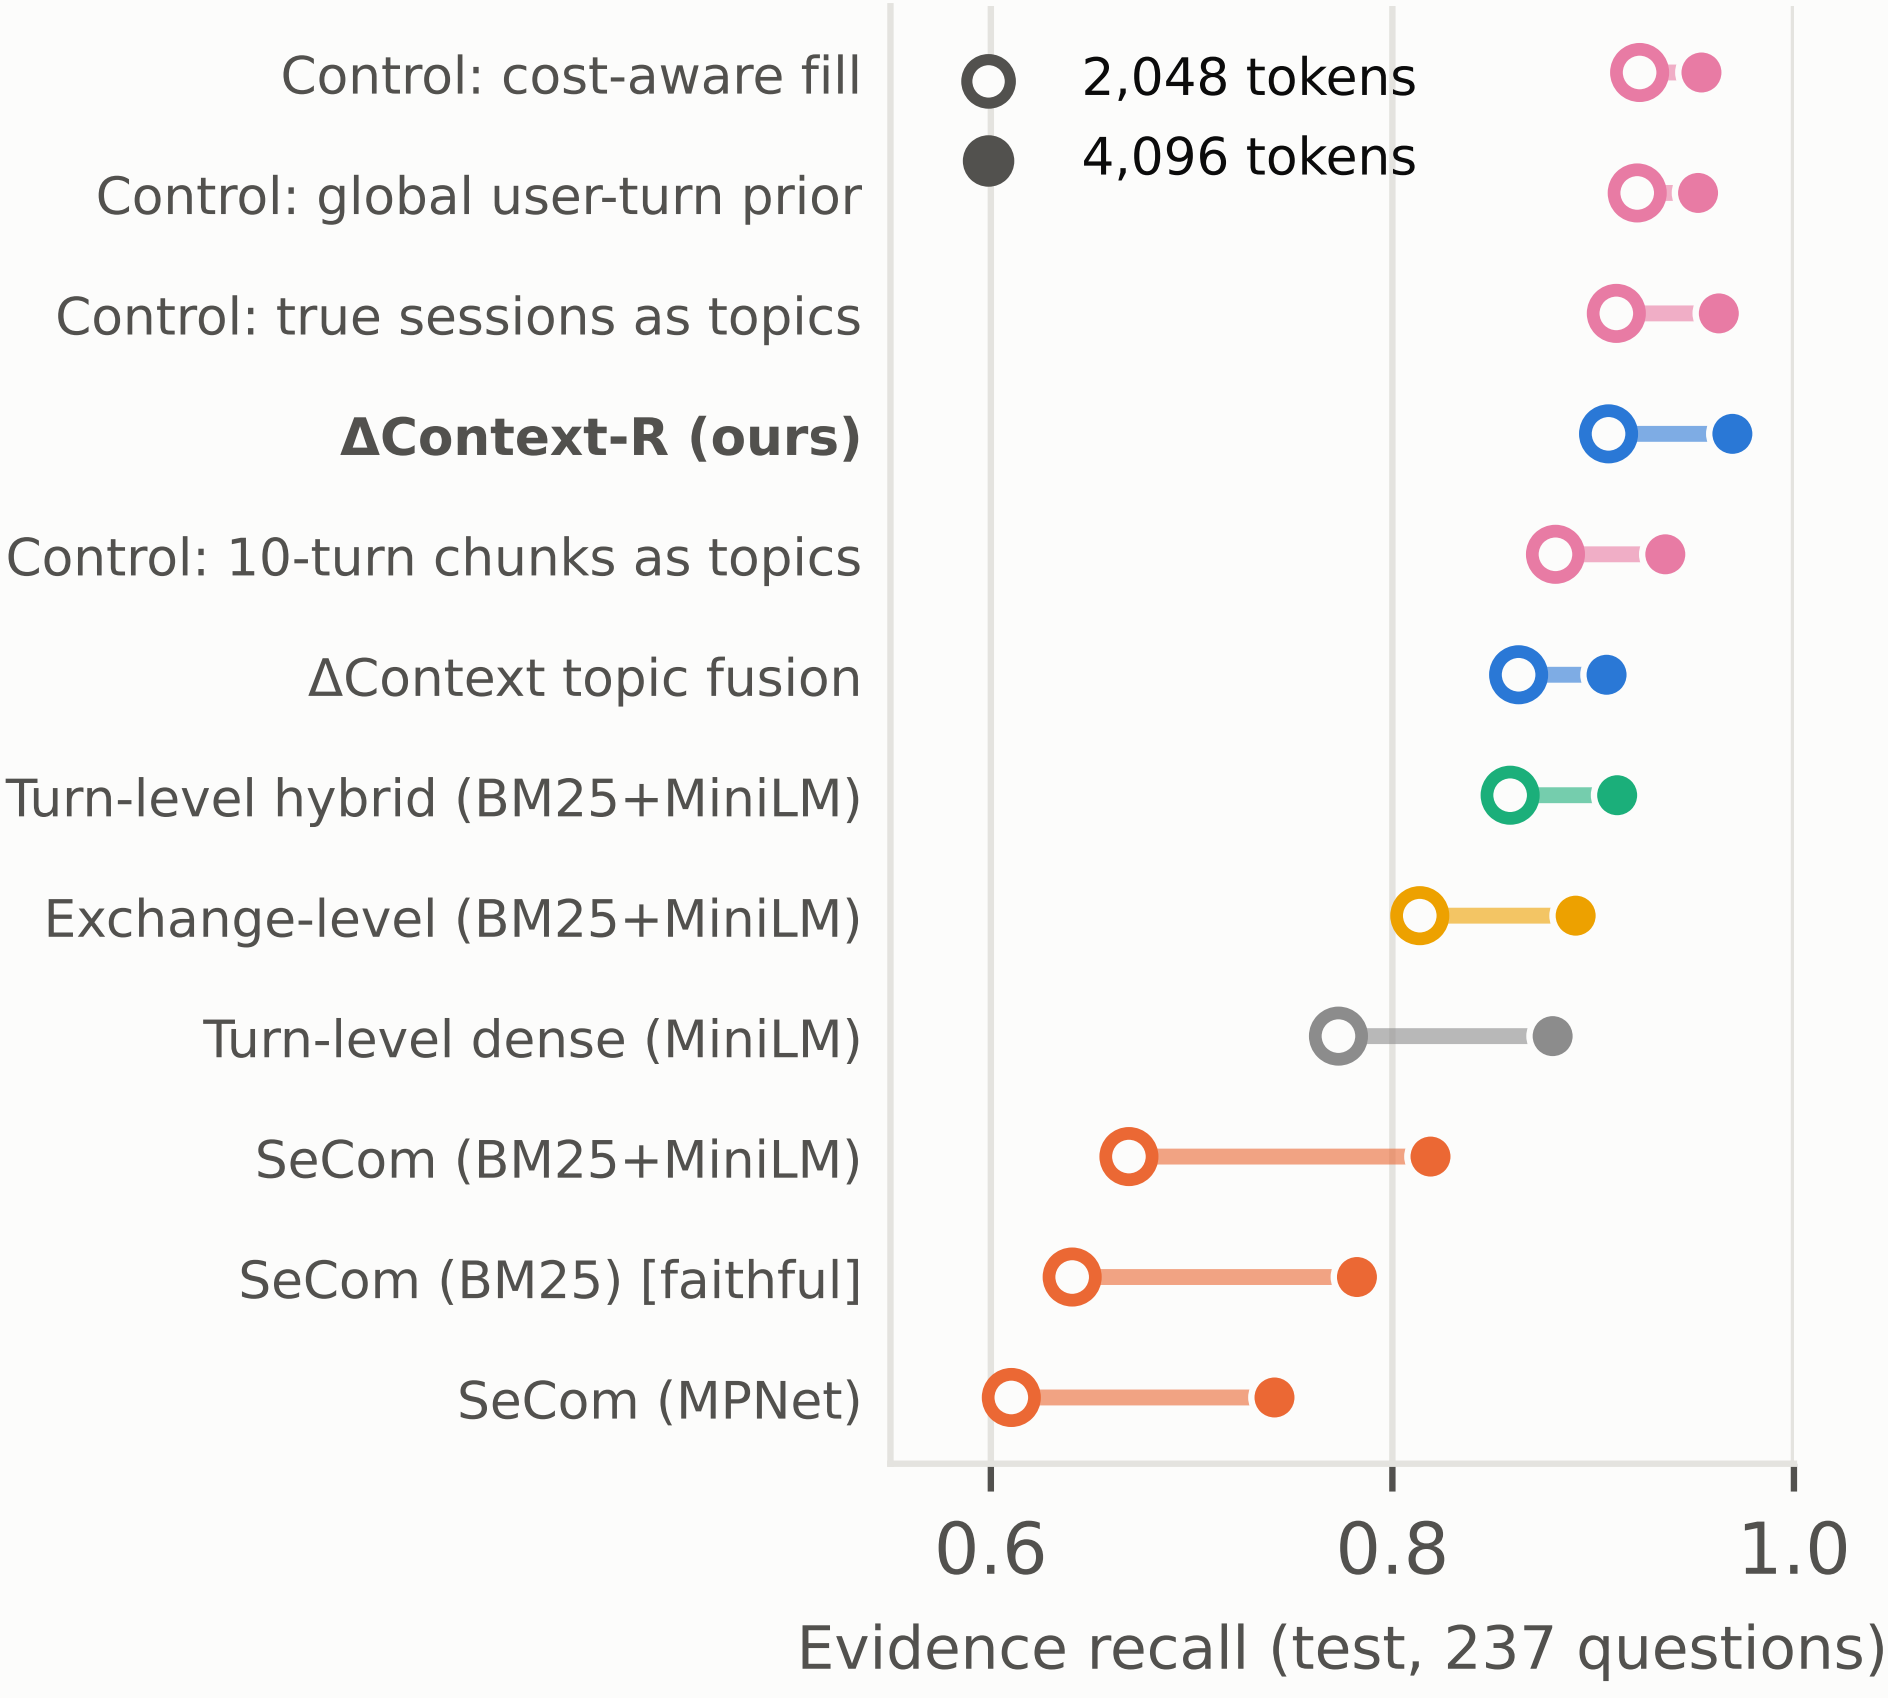
2,048 tokens
4,096 tokens
Control: cost-aware fill
Control: global user-turn prior
Control: true sessions as topics
ΔContext-R (ours)
Control: 10-turn chunks as topics
ΔContext topic fusion
Turn-level hybrid (BM25+MiniLM)
Exchange-level (BM25+MiniLM)
Turn-level dense (MiniLM)
SeCom (BM25+MiniLM)
SeCom (BM25) [faithful]
SeCom (MPNet)
0.6
0.8
1.0
Evidence recall (test, 237 questions)